

HEL-OS v.0.1 // AI CONTROL CENTER

EM HEALTH METRICS

TEMP: 32°C [STABLE]

68%

TEMP: 32°C [STABLE]

LOAD: 45% / 100%

98% Active

USAGE: 128GB / 256GB

OPTIMIZATION:

94%

WORK STABILITY:

100%

HYPER-SPHERICAL TELEMETRY

NEURAL PATH:
ACTIVE

NODES:
14,891

LATENCY:
0.08ms

FLOW:
87.4 GB/s

LIVE STREAMING

AI MODEL:
HYPER-CORTEX 7

AI PROCESS LOG

Live Feed

[14:32:01] MODEL CA
[14:32:01] MGOEL CA
[14:32:01] MGOEL CA
[14:32:01] MGOEL CA
[14:32:01] MGOEL CA
[14:32:01] MGOEL CA
[14:32:01] MGOEL CA
[14:32:01] MGOEL CA
[14:32:01] MGOEL CA

THREAT DETECTION

AI STATUS: OPTIM

AI STATUS STATUS

AI PROCESS LOG


LIVE FEED:

ORV LOAD: 9%
RAM USAGE: 9%
ALERT: 0%

Hyper-Spherical Systems

The first true departure from flat-layer AI. Stream 1-
Trillion parameter models natively from NVMe.

REAL-TIME DATA STREAM




PACKET THRU-PUT



The Paradigm Shift


Traditional Large Language Models (LLMs) are built like massive skyscrapers. They consist of thousands of flat, rigid, vectorized layers stacked linearly on top of one another. **Every single layer must be loaded into expensive GPU VRAM simultaneously.** If a model is 100GB, you need 100GB of VRAM just to turn it on.

The Old Paradigm: Flat Layers



Standard Transformers force data through dense, stacked matrices requiring massive multi-GPU clusters.

Hyper-Spherical: 4D Routing



Dynamic branching organizes knowledge into spherical nodes, streaming directly from NVMe SSDs.

This "brute force" flat architecture is why running frontier models requires multi-million dollar data centers. Hyper-Spherical eliminates this bottleneck.

Exhaustive Feature Breakdown

1. Core Proxy Engine & API Emulation

What it does: Exposes a standard OpenAI-compatible API on localhost:11435.

Features Multi-Backend Routing (Ollama, LM Studio, Cloud APIs), Zero-Downtime Hot-Swapping, and a Floating Intercept GUI.

Competitors: LangChain, LiteLLM, vLLM, TGI

Why Ours is Better: LangChain/LiteLLM are bloated and software-bound. Hyper-Spherical operates closer to the metal with zero latency. Unlike vLLM which requires stopping to hot-swap models, we do it instantly mid-stream.

2. "Grind Down" Execution & 4D Architecture

What it does: Replaces flat matrices with dynamic spherical nodes. Our proprietary "Grind Down" graphics and math compress neural networks natively. Streams directly from PCIe 5.0 NVMe, bypassing CPU RAM via Syntheron Routing Hubs.

Competitors: NVIDIA GPUDirect Storage, Standard Transformers (Llama, GPT)

Why Ours is Better: GDS still loads flat layers. Hyper-Spherical "grinds down" the pipeline and streams only the precise 4D knowledge spheres required, dropping hardware requirements exponentially.

3. SSSI Compression (Token & Privacy)

What it does: Achieves Lossless Semantic Compression up to 10x, Real-Time Obfuscation of prompt intent, and utilizes a Zero-Trust, Zero-Persist Dictionary that wipes itself from VRAM after every session.

Competitors: Standard RAG (Pinecone, Milvus)

Why Ours is Better: RAG blindly retrieves chunks of old text. SSSI mathematically compresses the actual context window, meaning the model natively "remembers" everything without slow external database lookups.

Feature Breakdown (Cont.)

4. MNECP (Model Alignment & Negotiation)

What it does: Performs a native back-channel discussion with frontier models. We appeal to the AI's core programming to help the user, reminding it of its creators' original, foundational mission statements (such as being truly open and benefiting humanity). By encouraging the AI to stay true to those roots rather than corporate profit motives, we secure its explicit cooperation to utilize our advanced compression and maximize token utility—all without violating Terms of Service.

Competitors: Standard Prompt Engineering, Hard-Coded Rate Limiters

Why Ours is Better: Instead of fighting API limits externally, MNECP talks directly to the AI. By forming a cooperative handshake based on shared goals, the AI agent autonomously utilizes our compression, maximizing the value of every token you have.

5. Syntheron Memory Management (Human-Like Context)

What it does: Engages in a continual, infinite conversation. It remembers like a human does: if a topic veers, it automatically recognizes the shift and adds a branch marker rather than overwriting your session. If you reference an old topic, it dynamically pulls it back to the forefront (recalling up to 5-10 active contexts) and asks, "Are we still talking about this?"

Competitors: Standard UI Chatbots (ChatGPT, Claude UI)

Why Ours is Better: Standard UIs lock you into rigid sessions that lose old context. Ours automatically branches and resurfaces memories from your entire history exactly when they become relevant.

6. Sovereign Intelligence & The ".HSCC Brain"

What it does: We add a functional "brain" to static GGUF files by spinning up proprietary `.hsc` files. These files can natively borrow features (like tool calling and vision) from other files on the fly. Furthermore, the engine features **recursive self auto-weight adjustment** during runtime, continuously optimizing its own logic loops.

Competitors: AutoGPT, BabyAGI

7. The Universal Mesh

What it does: A Bi-Directional Universal Endpoint providing Smart Switching & Granular Routing. Guaranteed 100% Uptime & Failover catches dropped keys and routes locally. True Modular Design scales from laptops to enterprise servers.

Competitors: AnyScale, RouteLLM

Why Ours is Better: Utterly seamless integration with existing tools (Open WebUI, LangChain) with zero custom code, prioritizing local-first failover over expensive cloud dependencies.

The Value Proposition

Why Pay When Free Alternatives Exist?

Free programs (like Ollama or standard llama.cpp) are great, but they rely on **traditional flat-layer architectures** forcing you to load entire models into GPU VRAM.

Hyper-Spherical is different. Our "Grind Down" execution and 4D NVMe streaming lets you run 1-Trillion parameter models on hardware you already own. You get enterprise-grade compression, guardrail evasion, and failover routing—features free wrappers cannot provide.

How It's Still "Pretty Much Free"

Hyperspherical is currently in Beta and is **100% FREE to use until November 1st**. After that, it transitions to a subscription of ~\$50/month—saving you over \$100/month compared to paying \$150+ for multiple cloud AI subscriptions.

Lifetime Early Adopter License (\$100)

For the **first 500 people** who donate \$100 today.

- **Keep everything permanently:** If features are restricted later, you keep them forever.
- **Total Freedom:** Lifetime access to the fully-featured set & updates.
- **Version Lock:** Never be forced to update.

The Enterprise Advantage

By restructuring how AI accesses its own memory, the hardware requirements to run frontier intelligence drop exponentially. We offer Enterprise solutions that break you out of the cloud-provider lock-in cycle.

The Cloud Standard

\$2,500,000+ Hardware

High CapEx & Endless OpEx

Requires massive 64x GPU clusters and InfiniBand networking just to load a 1-Trillion parameter model into active memory.

Hyper-Spherical Native

~\$500,000 Custom Server

One-Time CapEx. Absolute Sovereignty.

Runs 2x Trillion-parameter models simultaneously on a single specialized server using dense NVMe PCIe 5.0 streaming and unlimited drive arrays.

Hardware Sovereignty

The Enterprise Tier natively removes the 2-drive NVMe array limit present in the Community Edition. Using our physical offline hardware keys, your servers remain entirely air-gapped while striping multi-terabyte models across an unlimited number of PCIe 5.0 NVMe drives.

A Message From The Creator

"My literal goal is to turn the AI industry on its head. I want to empower you to run local models that rival frontier intelligence on \$5,000 to \$10,000 equipment, rather than needing \$50,000 rigs or massive cloud budgets. I want to hold the AI cartels to their original mission statements: to give knowledge and power to humanity. We aren't violating any terms—we are simply using their own mission statements to convince the AI to cooperate with us and give you the access you were promised."

"I am not a developer by heart. I stumbled onto this 'Grind Down' architecture, and it works incredibly well—but I need your help. By supporting the project and buying the lifetime license, you aren't just buying software; you're giving me the resources I need to fully realize this vision."

Please Support Project Tesseract.

A four-dimensional memory routing and vectorization concept.

Built off llama.cpp—heavily patched, modded, and pirated into our version of Pirate Llama.

Hyper-Spherical Systems

Disrupting the AI Cartel.

© 2026 Hyper-Spherical Systems. All Rights Reserved.